



PARALLEL COMPUTING

M.Sc. Embedded Systems



Report Part 1

Group Members:

Mehdi Moallemkolaei	153010947	mehdi.moallemkolaei@tuni.fi
Md Ashfak Haider Nehal	153045077	mdashfakhaider.nehal@tuni.fi

OCT 2025

Part 1.

Contents

Part 1.....	2
6.1. Benchmarking the original Code and Improving Performance via Compiler Settings....	3
6.1.1 & 6.1.2: Original code in Debug mode Vs. Release mode.....	3
6.1.3. Auto Vectorization Report:.....	3
6.1.4. Best Flags:.....	4
6.1.5. Each Flag's Performance.....	5
6.1.6. Code Breaker Flags:.....	5
6.2 Generic algorithm optimization.....	6
6.3. Multi-Thread Parallelization.....	7
6.3.1. loops are allowed to be parallelized?.....	7
6.3.2. Legal Loops but not beneficial?.....	7
6.3.3. Code Transformations.....	7
6.3.4. Transformation Effect.....	7
6.4. OpenMP Parallelization.....	8
6.4.1. Average Frame times.....	8
6.4.2. Extra code transformations/optimizations.....	8
6.4.3. Which loops were parallelized?.....	8
6.4.4. OpenMP break?.....	8
6.4.5. Scaling with core/hardware threads.....	8
6.4.6. Bonus Task.....	9
6.4.7. Machine and Environment.....	9
6.4.8. Hours to complete Part 1.....	9

6.1. Benchmarking the original Code and Improving Performance via Compiler Settings

6.1.1 & 6.1.2: Original code in Debug mode Vs. Release mode

It was tested on TC303 Lab PC:

Delays/Versions	Physics (moving sat)	Graphics (space coloring)	Total
1 Original C – Debug Mode	~ 170 ms	~ 1257 ms	~ 1429 ms
			~ 1461 ms
			~ 1407 ms
			~ 1399 ms
			~ 1433 ms
			~ 1465 ms
2 Original C – Release Mode	~ 119 ms	~ 560 ms	~ 681 ms
			~ 676 ms
			~ 693 ms
			~ 694 ms
			~ 670 ms
			~ 684 ms

The Release build improves performance by ~52% overall versus Debug, mainly due to compiler optimizations. First line in the table, shows the average over all frames after 10 first frames, and the other lines are sample frames randomly picked from 10 frames.

To run in the terminal (like for using seed):

```
.\out\build\x64-Release\parallel.exe
```

or

```
.\out\build\x64-Debug\parallel.exe
```

6.1.3. Auto Vectorization Report:

Did the compiler tell it managed to vectorize any of the loops in physics Engine or graphics Engine-functions?

No.

Enabling the flags in the CMake text file:

- `target_compile_options(parallel PRIVATE "/Qvec-report:2")`

The “Debug mode” Log does not contain any vectorization report.

The “Release mode” Log only shows:

- `info C5002: loop not vectorized due to reason 'xxxx'`

So, it means that no vectorization is done.

It reports like “C5002: loop not vectorized due to reason '1106/1203/1300/...'”.

For ParallelPhysicsEngine:

```
\Project\Satellites\parallel.c(129) : info C5002: loop not vectorized due to reason '1300'
\Project\Satellites\parallel.c(142) : info C5002: loop not vectorized due to reason '1203'
\Project\Satellites\parallel.c(138) : info C5002: loop not vectorized due to reason '1106'
\Project\Satellites\parallel.c(177) : info C5002: loop not vectorized due to reason '1300'
```

For ParallelGraphicsEngine:

```
\Project\Satellites\parallel.c(226) : info C5002: loop not vectorized due to reason '500'
\Project\Satellites\parallel.c(251) : info C5002: loop not vectorized due to reason '1104'
\Project\Satellites\parallel.c(196) : info C5002: loop not vectorized due to reason '1106'
```

6.1.4. Best Flags:

Experiment with the SIMD instruction set and FP relaxation-related optimization flags. Which compilation flags did you find to give the best performance?

We went through multiple scenarios to find the best flags:

Instructions used	Time spent on moving satellites	Time spent on space coloring	Total time in ms between frames
Ash PC			
/fp:fast	110 ms	574 ms	688 ms
/arch:AVX2	154 ms	983 ms	1142 ms
/fp:fast /arch:AVX2	97 ms	567 ms	668 ms
/fp:fast /arch:AVX	108 ms	567 ms	679 ms
/arch:AVX	153 ms	981 ms	1138 ms
/arch:sse4	154 ms	988 ms	1147 ms
/arch:avx512	153 ms	977 ms	1133 ms
/arch:sse4 /arch:avx512	153 ms	976 ms	1132 ms
/fp:fast /arch:avx512	110 ms	573 ms	686 ms
/fp:fast /arch:sse4	111 ms	574 ms	689 ms
/arch:AVX2 /arch:sse4	152 ms	984 ms	1141 ms
/arch:AVX2 /arch:avx512	154 ms	980 ms	1139 ms
Mehdi PC			
/fp:fast	87 ms	375 ms	465 ms
/arch:AVX2	131 ms	631 ms	765 ms
/fp:fast /arch:AVX2	86 ms	405 ms	493 ms
/fp:fast /arch:AVX	87 ms	355 ms	449 ms
/arch:AVX	131 ms	611 ms	744 ms
/arch:sse4	132 ms	615 ms	750 ms
/arch:avx512	127 ms	607 ms	736 ms
/arch:sse4 /arch:avx512	129 ms	610 ms	746 ms
/fp:fast /arch:avx512	86 ms	364 ms	458 ms
/fp:fast /arch:sse4	86 ms	364 ms	452 ms
/arch:AVX2 /arch:sse4	129 ms	602 ms	733 ms
/arch:AVX2 /arch:avx512	127 ms	601 ms	731 ms
Lab PC			
/fp:fast	83 ms	334 ms	420 ms
/arch:AVX2	118 ms	548 ms	668 ms
/fp:fast /arch:AVX2	81 ms	368 ms	450 ms
/fp:fast /arch:AVX	78 ms	324 ms	404 ms
/arch:AVX	122 ms	561 ms	685 ms
/arch:sse4	121 ms	548 ms	670 ms
/arch:avx512	118 ms	555 ms	676 ms
/arch:sse4 /arch:avx512	122 ms	556 ms	680 ms
/fp:fast /arch:avx512	78 ms	331 ms	411 ms
/fp:fast /arch:sse4	81 ms	331 ms	413 ms
/arch:AVX2 /arch:sse4	122 ms	547 ms	671 ms
/arch:AVX2 /arch:avx512	123 ms	545 ms	669 ms

flags “/fp:fast /arch:AVX” gave the best results on 2/3 of tested PC setups. So, we choose this combo for the rest of our project.

However, this combo “/fp:fast /arch:AVX2” also showed a great result and we will also use this combo to compare with the first chosen combo to see which will give the best final result.

* Update: both combinations are tested on the final code and they both have the exact same final result.

6.1.5. Each Flag's Performance

Can you explain what each of the optimization flags you found to give the best performance does?

- `/fp:fast` - Relaxes strict IEEE-754 rules so the compiler can reorder and "contract" FP ops (e.g., form FMAs), use approximate `recip/rsqrt`, assume no NaNs/Infs, and reassociate expressions. This often unlocks SIMD for `sqrt/div` heavy code and reduces dependencies. Result: faster but slightly different numerics.
- `/arch:AVX` - Allows the autovectorizer to target AVX (256-bit YMM registers) for float/double SIMD. Wider vectors than SSE, but no integer AVX2 ops.
- `/arch:AVX2` - Enables AVX2 (still 256-bit, but with rich integer/vector ops and gathers). On FP workloads, also lets the compiler use FMA3 where profitable when contraction is allowed (e.g., with `/fp:fast`). Often the sweet spot on modern CPUs.
- `/arch:AVX512` - Permits 512-bit vectors, mask registers, and wide loads/stores. Only used if the CPU/OS support it; many consumer CPUs don't, or the compiler may avoid it due to potential down-clocking. If unsupported, the compiler falls back to a lower ISA.
- `/arch:SSE4` - Targets 128-bit SSE4.x. Narrower vectors; useful for older CPUs, typically slower than AVX/AVX2 on modern chips.

6.1.6. Code Breaker Flags:

Did you find some compiler flags which cause broken code to be generated, and if so, can you think why?

No breakage occurred.

6.2 Generic algorithm optimization

Can you find any ways to change the code to either get rid of unnecessary calculations, or allow the compiler to vectorize it better, to make it faster? If yes, what did you do and what is the performance with your optimized version?

1. Inside parallelGraphicsEngine(), replacing the sqrt and merging two satellite loops

Why it's faster

- Removes 2× sqrtf per pixel (black hole & first satellite loop).
- Eliminates the second satellite loop entirely (we accumulate numerator+weights on the first pass and normalize once).
- Replaces divisions by using the reciprocal once (invW) and squared-distance comparisons (SIMD-friendly).
- Keeps byte-for-byte semantics vs your current sequential engine (we preserved the “nearest baseline + 3-weighted average” behavior and the inside-radius early white).

Vectorization angle

- Squared distance + branch consolidation turns the inner loop into straight-line, division-light code. Compilers are far more willing to autovectorize $dx \cdot dx + dy \cdot dy$, reciprocal, FMAs, and min-tracking than a path with sqrt + break + two passes.

2. Inside parallelGraphicsEngine(), removing % and /

Why is it faster

- Eliminates one integer division and one modulo per pixel (both costly).
- Access patterns become perfectly linear, which the compiler and CPU prefetcher love.

Vectorization angle

Regular 2D loops + linear idx often trigger better auto vectorization and improved cache behavior.

Performance:

		moving satellites	space coloring	Total time
Ash PC	Before	97 ms	567 ms	668 ms
	After	97 ms	544 ms	643 ms
Mehdi PC	Before	87 ms	355 ms	449 ms
	After	85 ms	354 ms	444 ms
Lab PC	Before	78 ms	324 ms	404 ms
	After	76 ms	321 ms	400 ms

6.3. Multi-Thread Parallelization

6.3.1. loops are allowed to be parallelized?

Which of the loops are allowed to be parallelized to multiple threads?

No.	Loop	Legal?	Worth doing?	Reasoning
1	Physics iteration	No	–	Loop-carried time dependence (step t feeds step t+1).
2	Physics satellite	Yes	Yes	Independent per-satellite within a time step.
3	Graphics pixel	Yes	Definitely!	Massive, independent work per pixel. Ideal scaling.
4	Graphics satellite	Yes	Almost No	The trip count is small. Expensive to thread. Prefer SIMD inside each pixel.

6.3.2. Legal Loops but not beneficial?

Are there loops which are allowed to be parallelized to multiple threads, but which do not benefit from parallelization to multiple threads? If yes, which and why?

Physics satellite – Allowed, but not beneficial: putting `#pragma omp for` inside the time-step loop causes PHYSICSUPDATESPERFRAME (=100 000) barriers per frame; overhead kills scaling.

Graphics satellite – Allowed, but not beneficial: only 64 iterations, branchy (break) and reduction-like pattern → threads add overhead; better target is SIMD.

6.3.3. Code Transformations

Can you transform the code in some way (change the code without affecting the end results) which either allows parallelization of a loop which originally was not parallelizable, or makes a loop which originally was not initially beneficial to parallelize with OpenMP beneficial to parallelize with OpenMP? If yes, explain your code transformation?

Transformation 1 – Physics: loop interchange + register temporaries

What changed: Swapped the loop order so each thread owns a block of satellites and advances each one through all-time steps; kept `x,y,vx,vy` in registers and wrote back once at the end; used `schedule(static)`.

Why it helps: Removes the 100,000 per-step barriers and avoids false sharing from repeatedly writing adjacent `tmpPosition[i]/tmpVelocity[i]`. Coarse grains per thread ⇒ good scaling.

Side math cleanup: Hoisted `dt = DELTATIME/PHYSICSUPDATESPERFRAME`, replaced two divides by `|r|` with one sqrt plus multiplies (`invd, invd2`).

Transformation 2 – Graphics: parallelize rows, keep one-pass inner loop (no sqrt)

What changed: Parallelized the Graphics pixel loop by rows (`#pragma omp parallel for schedule(static) over y`); turned the pixel loop into a 2-D sweep (no `%//` per pixel); fused the two satellite passes into one pass that (a) early-outs to white if inside radius, (b) accumulates the weighted sums and weights, and (c) tracks the nearest using squared distance; normalized once.

Why it helps: Big, cache-friendly chunks (rows) minimize scheduling overhead and keep writes linear; removing `sqrtf` and extra pass slashes scalar cost and makes the inner loop SIMD-friendlier.

6.3.4. Transformation Effect

Does your code transformation have any effect on vectorization performed by the compiler?

Yes, replacing `sqrtf/two-pass` logic with squared-distance + one pass reduces divisions and branches, which helps the compiler's autovectorizer (and at least improves instruction-level parallelism even if SIMD is limited). The physics inner loop still has a time dependency, so SIMD across steps is not expected; the cleaned math reduces stalls.

6.4. OpenMP Parallelization

6.4.1. Average Frame times

What are the average frametimes (milliseconds) in your OpenMP-multithreaded version of the code?

Setup used		Time spent on moving satellites	Time spent on space coloring	Total time in milliseconds between frames
Own PC Ash	Without OpenMP	97 ms	544 ms	643 ms
	With OpenMP	11 ms	84 ms	97 ms
Own PC Mehdi	Without OpenMP	95 ms	195 ms	293 ms
	With OpenMP	8 ms	24 ms	34 ms
Lab PC	Without OpenMP	87 ms	170 ms	261 ms
	With OpenMP	5 ms	16 ms	23 ms

6.4.2. Extra code transformations/optimizations

Physics (T1): Loop interchange so that each thread owns a subset of satellites and advances them through all time steps (removes 100k step barriers). Kept x, y, v_x, v_y in registers and wrote back once to avoid false sharing. Hoisted $dt = \text{DELTA_TIME} / \text{PHYSICS_UPDATES_PER_FRAME}$ and replaced two divides-by- $\|r\|$ with one sqrt + multiplies.

Graphics (T2): Parallelized rows (outer pixel loop), changed to a 2-D sweep (no per-pixel $\%/\%$), removed sqrtf (squared-distance tests), and fused the two satellite passes into one (accumulate weights & nearest in one loop, then normalize once).

6.4.3. Which loops were parallelized?

Physics satellite (inner): Yes, after loop interchange (OpenMP over satellites, one write-back per satellite).

Graphics pixel (outer): Yes, OpenMP over rows (schedule(static)).

6.4.4. OpenMP break?

Did any OpenMP parallelization break or slow things down?

- Naïve Physics satellite inside the time-step loop: Slowed down due to ~100,000 barriers per frame; also, more false sharing from per-step writes to adjacent `tmpPosition[i]/tmpVelocity[i]`. Fix: loop interchange + keep per-satellite state in registers, write back once, `schedule(static)`.
- Graphics pixel with collapse (2) and aggressive SIMD on MSVC: Slight regression (scheduling overhead + conservative SIMD); the row-parallel version with one-pass inner loop was consistently faster.

6.4.5. Scaling with core/hardware threads

Did the performance scale with the number of CPU cores or native CPU threads (if you have an AMD Ryzen, or Intel Core i7 or i3 CPU, cores may be multi-threaded and can execute two threads simultaneously)? If not, why?

Performance improved substantially when OpenMP was enabled and loops were coarse-grained (rows/satellites). On typical 6-16 logical-thread CPUs, Graphics pixel dominates scaling; Physics scales well after loop-interchange. If scaling plateaus, the reasons are:

- I. Graphics-bound: once graphics time dominates, adding threads to physics yields little change.
- II. Runtime overhead: too-fine task slicing (fixed by `schedule(static)` and coarse-grained).
- III. Memory bandwidth/cache effects: especially at high thread counts.

6.4.6. Bonus Task

As a bonus question, you are now free to implement any CPU optimizations (incl. advanced multicore and vectorization optimizations) to your program as long as your algorithm produces correct results and passes the error check. How fast can you get your program on a CPU? Please, report your methods.

The additional work on structures and mathematics to reduce latency is described in each related question in this report. like in [6.3.3](#), [6.4.2](#), or other answers.

6.4.7. Machine and Environment

We did the project both on a computer at TC303 and on our own PC to test the functionality and differences.

Setup	Configuration
Ashfak PC	Processor: Intel(R) Core (TM) i5-10400F CPU @ 2.90 GHz RAM: 16GB Compiler: Visual Studio x64 Version: 4.8.09032
Mehdi PC	Processor: Intel(R) Core (TM) i7-12700H CPU @ 2.30 GHz RAM: 16GB Compiler: Visual Studio x64 Version: 4.8.09032
Lab PC	Processor: Intel(R) Core (TM) i7-13700 CPU @ 2.10 GHz RAM: 64GB Compiler: Visual Studio x64 Version: 4.8.09032

6.4.8. Hours to complete Part 1

We spent approximately two weeks finishing part 1.
~2 hr/day, total of around 30 hrs.